

The Non-Relational Blind Spot: Evaluating LLMs on Multi-Table Numerical Reasoning

Anonymous ACL submission

Abstract

Large language models (LLMs) are increasingly used for tabular question answering, yet most evaluations assume clean, relational schemas. This leaves a gap for real-world analytical queries, where evidence may be distributed across multiple non-relational tables and encoded through layout, noise, units, or implicit cross-table links. Existing relational datasets are difficult to adapt to this setting because complex perturbations can break table validity and answer verification. We introduce GRADINO, a from-scratch benchmark generator for controlled numerical QA evaluation over multiple non-relational tables. GRADINO synthesizes a latent relational source for executable SQL supervision, then evaluates models on perturbed non-relational views derived from the same source. This separation preserves answer correctness while controlling reasoning topology, table count, layout perturbations, and unit heterogeneity. Our evaluation of LLMs across financial and environmental domains reveals a critical blind spot: non-relational formatting, unit discrepancies, and multi-table scaling cause sharp, compounded performance drops.

1 Introduction

Large language models (LLMs) are increasingly used as natural-language (NL) interfaces for querying tabular data, making Table Question Answering (Table QA) a natural testbed for evidence retrieval, numerical operations, and analytical reasoning over structured information. Most Table QA benchmarks assume that structure is *explicit*, as in clean relational schemas, and that the relevant evidence is contained in a single table or connected through predefined database relations. Many real spreadsheets and reports are instead *non-relational*: the structure needed to answer a question is encoded in the presentation itself, through hierarchical headers, merged cells, blank regions, unit conventions,




Figure 1: Motivating comparison between relational and non-relational presentations: each question is evaluated on relational and non-relational variants of the same underlying data in both single- and multi-table settings. The figure reports exact-match accuracy (%) for gpt-5-mini and Qwen-family models on a sample of 192 questions from Spider and Kaggle datasets.

noisy entries, or implicit dependencies within and across tables (Cheng et al., 2022).

This is a blind spot in current Table QA evaluation, as highlighted by the W3C 'Data on the Web Best Practices' report(?). A model may answer correctly over a clean relational table yet fail on a non-relational rendering of the same evidence, especially when that evidence is distributed across multiple tables (Figure 1; §C). Characterizing this failure mode requires controlled benchmarks that preserve the underlying evidence and gold answer across alternative renderings, while allowing systematic variation in domain, layout, perturbations, unit conventions, and cross-table evidence distribution.

Existing Table QA resources address only part of this need and fall into two families. (i) *Manually curated* relational and non-relational benchmarks offer realistic tables (§2) but fix the domain, layout, question distribution, reasoning regime, and table count, with few targeting genuine multi-table reasoning (Zhao et al., 2022; Contalbo et al., 2025). Extending them means hand-crafting new tables, questions, units, missing values, perturbations, and

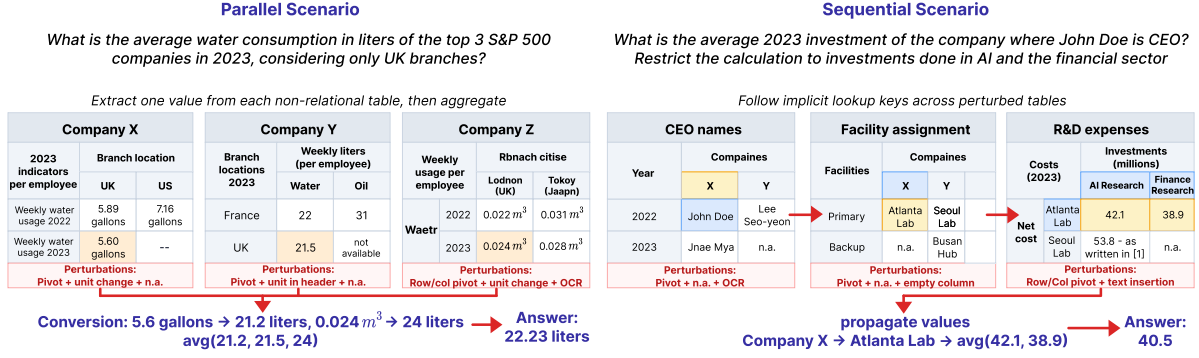


Figure 2: Multi-table parallel and sequential scenarios generated by GRADINO over non-relational tables. Parallel questions extract and aggregate values across perturbed tables; sequential questions propagate lookup keys across tables before aggregating final values.

gold answers, which is costly to scale, hard to adapt, and increasingly contaminated in recent LLMs (Ranaldi et al., 2024). (ii) *Automatic generation* methods are scalable and yield verifiable supervision, but operate over existing relational databases (Papicchio et al., 2023; Oh et al., 2024; Park et al., 2026), which dictate what can be generated: the practitioner cannot freely set the domain, perturbations, unit conventions, or the number of tables a question spans, precisely the axes needed to probe non-relational multi-table reasoning.

To probe these axes and generate non-relational Table QA instances, one natural workaround would be to convert relational tables from existing Table QA benchmarks into non-relational views. However, post-hoc conversion falls short on two counts. (i) *Practically*, most relational tables resist dense hierarchical rendering: pivoting their attributes yields sparse layouts or missing combinations, so only a small fraction of Spider (Yu et al., 2018), BIRD (Li et al., 2023), and BEAVER (Chen et al., 2024) tables convert into realistic dense or hierarchical views, and this constraint becomes even tighter in multi-table scenarios (§C). (ii) *By design*, converted data inherits the source’s schema, content, and relations, granting no independent control over reasoning structure, evidence distribution, and tabular presentation. This lack of control also limits question diversity, as non-relational multi-table QA may contain two complementary evidence-flow question patterns (Figure 2): *parallel*, where values retrieved independently from different tables are aggregated, and *sequential*, where intermediate lookup keys are propagated across tables. A non-relational Table QA generator must therefore build novel benchmarks from scratch, which is the only way to produce instances diverse in domain,

perturbations, table count, and question patterns.

We introduce GRADINO, a from-scratch framework for generating automatically verifiable QA benchmarks over multiple non-relational tables, with explicit control over parallel and sequential evidence-flow patterns. Unlike post-hoc conversion, GRADINO synthesizes its own latent source rather than adapting an existing database. Given a target domain and lightweight controls such as attribute cardinalities and table count, GRADINO uses an LLM to define a domain-specific schema and instantiate a *latent relational source*. From this source, it derives executable SQL programs with scalar answers, renders perturbed *non-relational views*, and verbalizes the programs into NL questions. The rendering stage combines established Table QA perturbations (Gan et al., 2021; Pi et al., 2022; Chang et al., 2023; Zhang et al., 2025a) with pivots, hierarchical headers, merged cells, and unit conversions. This design *decouples supervision from presentation*: answers are computed and verified on the latent relational source, while tested models see only the perturbed non-relational views.

We evaluate GRADINO on financial and environmental scenarios with up to 20 parallel and 5 sequentially connected tables per question, using gpt-5-mini and Qwen-family models. As table count increases, accuracy drops by up to 39.6 points in parallel settings and 32.5 points in sequential settings, while unit heterogeneity causes additional losses of up to 50.1 points, with strong domain dependence. The drop generalizes to real tables: on the few real tables admitting non-relational rendering, perturbation costs 7.95 points single-table and 30.05 multi-table (§C). Manual inspection finds 95.9% of 240 generated questions correct. Finally, GRADINO-generated data comple-

Class	Approach	Generated artifacts			Reasoning / evaluation control			
		Tbl.	NLQ	Ans.	Pert.	Rel.	Non-rel.	# tables
Perturbation	Dr.Spider (Chang et al., 2023)	✗	✗	✗	Sch/Txt/Cell	✓	✗	SB
	EvoSchema (Zhang et al., 2025a)	✗	✗	✗	Sch	✓	✗	SB
	ADVETA (Pi et al., 2022)	✗	✗	✗	Sch	✓	✗	SB
	FREB-TQA (Zhou et al., 2024)	✗	✗	✗	Cell/Lay	✗	✓	1
Tabular synthesis	REaLTabFormer (Solatorio and Dupriez, 2023)	✓	✗	✗	✗	✓	✗	N/A
	GOGGLE (Liu et al., 2023)	✓	✗	✗	✗	✓	✗	N/A
	TabGen-ICL (Fang et al., 2025)	✓	✗	✗	✗	✓	✗	N/A
QA / SQL synthesis	Liu et al. (2022)	✗	✗	✓	✗	✓	✗	1
	Jiang et al. (2022)	✗	✓	✓	✗	✓	✗	1
	Pal et al. (2023)	✗	✗	✓	✗	✓	✗	SB
	ER-Bench (Oh et al., 2024)	✗	✓	✓	✗	✓	✗	SB
	SPARTA (Park et al., 2026)	✗	✓	✓	✗	✓	✗	SB
	DSQG-Syn (Duan et al., 2025)	✗	✓	✓	✗	✓	✗	SB
	SQLForge (Guo et al., 2025)	✓	✓	✓	✗	✓	✗	SB
	SYNQL (Baumgartner and Kornuta, 2024)	✗	✓	✓	✗	✓	✗	SB
Ours	GRADINO	✓	✓	✓	Sch/Cell/Lay/Num	✓	✓	GC

Table 1: Comparison with data-generation approaches. Tbl., NLQ, and Ans. denote generated tables, natural-language questions, and automatically verifiable answers. Perturbations are abbreviated as schema/name (Sch), text/query (Txt), cell/content (Cell), layout/structure (Lay), and numerical/unit (Num). The number of tables is either fixed to one, source-bounded (SB), or generator-controlled (GC).

ments scarce in-domain supervision, as combining GRADINO with the small GRI-QA (Contalbo et al., 2025) training set yields the best results on the GRI-QA test set.

In summary, our contributions are: (i) we introduce GRADINO, a from-scratch generator of automatically verifiable numerical QA over multiple non-relational tables; (ii) we define controllable parallel and sequential evidence-flow patterns with answer-preserving structural, cell-level, and unit-based perturbations; (iii) we show that table count, non-relational presentation, and unit heterogeneity substantially degrade LLM performance, and confirm on perturbed real tables that non-relationality is especially harmful in multi-table settings; (iv) we validate generated-question correctness and show that GRADINO-generated data provides effective supervision that complements scarce in-domain data; (v) we release code and generated datasets¹.

2 Related Work

Table QA benchmarks and robustness to tabular perturbations. Early Table QA benchmarks focused on relational tables, from single-table semantic parsing (Zhong et al., 2017) to multi-table relational QA (Yu et al., 2018; Wu et al., 2025a; Li

et al., 2023; Chen et al., 2024). Subsequent studies showed that these systems remain brittle under schema, query, cell, and layout perturbations (Gan et al., 2021; Pi et al., 2022; Chang et al., 2023; Zhang et al., 2025a; Zhou et al., 2024; Liu et al., 2024), as well as in realistic or domain-specific sources (Furst et al., 2025; Qiu et al., 2024).

Non-relational Table QA benchmarks instead target semi-structured, hierarchical, or hybrid table-text inputs, where merged headers, layout cues, and irregular cell links make direct Text-to-SQL execution unsuitable (Pasupat and Liang, 2015; Herzig et al., 2021; Chen et al., 2020a; Nan et al., 2022; Cheng et al., 2022; Wu et al., 2025b; Chen et al., 2020b, 2021a). Beyond structural complexity, domain-specific benchmarks introduce specialized terminology and numerical reasoning challenges across areas such as finance, science, and environmental reporting (Zhu et al., 2021; Chen et al., 2021b, 2022; Deng et al., 2022; Reddy et al., 2024; Katsis et al., 2022; Lu et al., 2023; Zhang et al., 2025b; Ghosh et al., 2024). Moreover, only a few benchmarks address non-relational multi-table reasoning: MultiHiertt (Zhao et al., 2022) targets sequential reasoning over financial tables and text, whereas GRI-QA (Contalbo et al., 2025) requires parallel aggregation across environmental reports.

Although valuable, manually curated bench-

¹Anonymous Github link

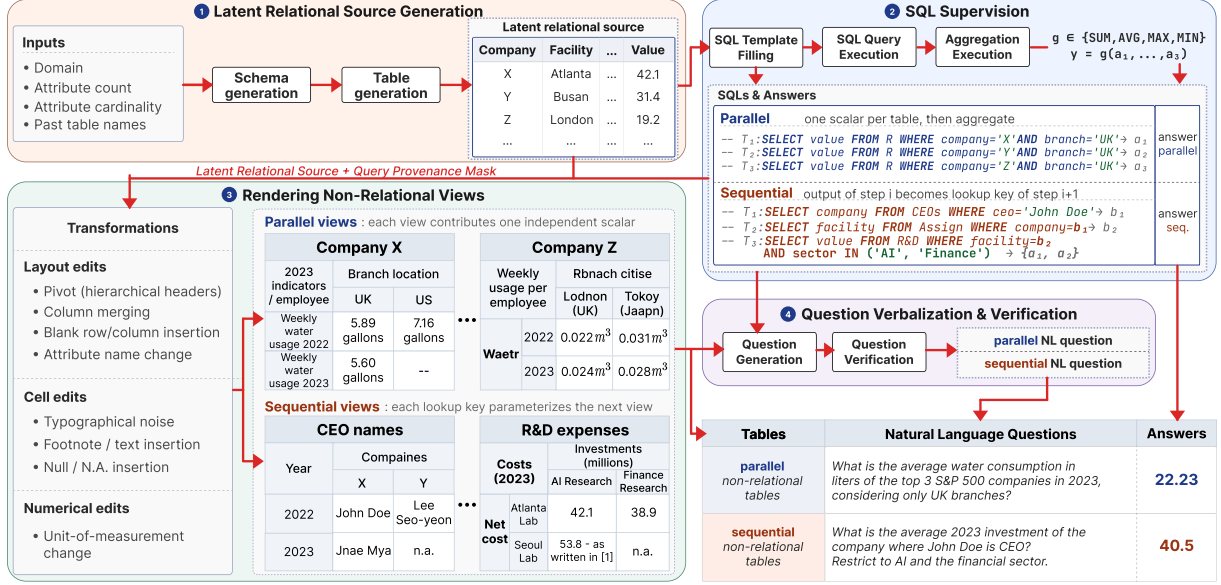


Figure 3: Overview of GRADINO. A latent relational source provides SQL supervision and gold answers, while models are evaluated only on perturbed non-relational views.

marks are largely fixed in their domains, layouts, question distributions, perturbations, reasoning regimes, and number of tables. In contrast, GRADINO is a configurable generation framework that provides control over parallel and sequential reasoning, table count, and structural, cell-level, layout, and numerical perturbations. Compared with existing benchmarks, this configurability enables systematic evaluation, model auditing, and domain-specific multi-table training at a scale that manual curation cannot support (§4.4).

Automatic generation of tabular data. Synthetic Table QA and Text-to-SQL generation reduce annotation cost by sampling executable programs and verbalizing them into natural language, using SQL templates, query-diversification methods, or relational multi-hop constructions (Jiang et al., 2022; Pal et al., 2023; Baumgartner and Kornuta, 2024; Guo et al., 2025; Duan et al., 2025; Oh et al., 2024; Park et al., 2026). These methods usually start from existing relational databases, so they inherit fixed schemas, table counts, and question spaces. A separate line generates synthetic tables from examples or learned relational structure (Fang et al., 2025; Solatorio and Dupriez, 2023; Liu et al., 2023; Hoppe et al., 2025), but targets tabular synthesis rather than QA. As summarized in Table 1, existing approaches generally address perturbation, table generation, or verifiable QA synthesis separately. In contrast, GRADINO generates latent sources, non-relational views, executable answers, and multi-table questions from scratch, while al-

lowing user guidance when higher domain realism is needed (Ahmad et al., 2025) (§B).

3 Approach

GRADINO follows the pipeline shown in Figure 3. Given a domain specification and a small set of generation controls, GRADINO constructs benchmark instances in four stages: (i) it generates a *latent relational source* (§3.1); (ii) it applies executable SQL programs and computes their gold answers (§3.2); (iii) it renders the latent source as one or more perturbed *non-relational views*, preserving answer correctness through provenance masks (§3.3); and (iv) it verbalizes the SQL programs into NL questions (§3.4). The generated NL questions and non-relational views are used for model evaluation.

3.1 Latent Relational Source Generation

This stage constructs the latent relational source in two steps. First, GRADINO prompts an LLM to generate a domain-specific schema, conditioned on a user-defined target domain, attribute count, and attribute cardinalities, together with the names of previously generated tables, which are supplied automatically by the system from prior generation steps. This encourages related but non-identical tables. The schema specifies the table name; the attributes, each with its type and a corresponding value range; unit metadata; and a designated numerical attribute for quantitative reasoning. To steer generation toward realistic domains, the prompt can be grounded in example tables and domain-specific unit inven-

ories, such as the Quantities, Units, Dimensions and Types ontology (QUDT) (qud) for environmental data and the XBRL Units Registry (XBRL UTR) (XBRL International) for financial data.

Second, GRADINO instantiates the schema as a dense latent relational source by enumerating combinations of categorical attributes and sampling the numerical field. During instantiation, it enforces lightweight semantic regularities, including column dependencies (e.g., city $\rightarrow$ country) and cross-row numerical constraints (e.g., a “gross” value exceeds its corresponding “net” value for the same entity and period). We defer their formalization and solver implementation to §B.

The resulting latent relational source serves as the unified basis for SQL supervision, question generation, and the construction of multiple non-relational views. Rather than generating these views independently, GRADINO derives them from the same latent representation through table-specific projections and transformations.

3.2 SQL Supervision and Multi-Table Reasoning

This stage uses SQL as supervision: GRADINO executes SQL queries over the latent relational source to compute gold labels. We focus on aggregation-style questions that combine scalar values.

In the *parallel* setting, GRADINO constructs a predicate over categorical attributes that isolates a target numerical value. It instantiates multiple table-specific selections by keeping the predicate fixed except for one attribute that varies across views (e.g., keeping the sector fixed while changing the company; see Figure 3). Each view contributes one scalar a_i , and the final answer is computed as $y = g(a_1, \dots, a_n)$, with $g \in \{\text{sum, avg, max, min}\}$. Thus, parallel instances correspond to multi-table *quantitative* (sum, average), and *comparative* questions (superlative operations like max, min), with one extractive SQL query per view and one aggregation over the scalars.

In the *sequential* setting, GRADINO constructs a chain of table-local queries. Intermediate views are associated with extractive queries whose returned value becomes the only lookup key for the next view in the chain, while the final view applies the target numerical aggregation. The final aggregation uses the same operators as in the parallel setting, applied to a randomly sampled set of 2 to 5 values returned by the last view. The gold label is the result of executing this latent composition. This setup

yields multi-step reasoning across tables without relying on an externally given relational schema. For instance, the example in Figure 3 requires identifying the company associated with John Doe, using that company to retrieve the AI and financial-sector investment values, and averaging them.

Since the same latent source can be reused to generate multiple types of questions and different numbers of views, the approach scales naturally between reasoning depths and evidence cardinalities.

3.3 Rendering Non-Relational Views

Rendering is organized as an answer-preserving transformation sequence. GRADINO first creates table-specific relational projections from the latent source and records a provenance mask marking cells required by the SQL program and cells safe to perturb. It then introduces null values, and applies column merging, unit rendering and conversion while the representation is still relational. Next, GRADINO converts each projection into a non-relational view through a pivot-style transformation, organizing selected attributes as hierarchical row and column headers. The generator searches over candidate row/column partitions, prefers dense pivots, shrinks the layout, and retains all evidence required by the SQL program. Rendered views are capped at a user-specified size. After layout construction, GRADINO inserts perturbations such as blank rows or columns, typographical noise, and contextual or footnote-like text.

The same rendering machinery supports both multi-table regimes, with scenario-specific controls. In parallel generation, each view provides one independently retrievable scalar for the final aggregation, and unit conversion can force normalization across compatible quantities. In sequential generation, each view represents one step of a latent lookup chain. Intermediate views contain exactly one occurrence of the lookup key extracted from the previous view. This is an evaluation choice rather than a framework limitation: ambiguous or tree-shaped joins would add uncontrolled variation in join complexity and the number of values entering the final aggregation. Unique matches keep the reasoning depth comparable across instances while still requiring models to recover implicit cross-table links, since latent joins are not stated in the question. Finally, view order is randomized to prevent shortcutting. Overall, rendering hides clean schemas, explicit joins, and standardized numerical expressions while preserving verifiability through




Figure 4: Scaling behavior as the number of tables increases. Labels at the end of curves in (a) and (b) report the accuracy drop from the smallest to largest table count. Panel (c) compares parallel and sequential questions at matched table counts, showing that sequential reasoning adds degradation beyond multi-table aggregation.

the latent relational source and SQL supervision.

3.4 Question Verbalization and Verification

In the final stage, GRADINO prompts an LLM to generate a natural-language question. The prompt receives all non-relational views in HTML form, the SQL supervision, the aggregation description, and the gold answer. The multi-table verbalization differs slightly across the two reasoning regimes. For parallel questions, the generator verbalizes the shared selection pattern together with the cross-table aggregation to be performed. For sequential questions, the generator verbalizes only the first and last steps of the chain and leaves intermediate dependencies implicit in the table contents. As a result, the final question does not expose the latent composition directly, but still requires the evaluated LLM to recover it from the views. To filter out noisy generations, we perform an LLM-based verification pass after question generation, repairing the NL question whenever the verbalization is inconsistent with the underlying supervision.

4 Experimental evaluation

Our experiments are designed to answer four questions. **(RQ1) Scaling:** how does performance evolve as the number of tables grows, in both parallel and sequential settings? **(RQ2) Numerical heterogeneity:** to what extent do unit discrepancies and quantitative computation degrade accuracy, and is this effect domain-dependent? **(RQ3) Utility:** beyond evaluation, can GRADINO-generated data serve as effective training supervision? **(RQ4) Question correctness:** what is GRADINO’s percentage of correctly generated questions? We address these through controlled benchmarks in the finan-

cial and environmental domains, and validate our findings on perturbed real tables in §C.

4.1 Baselines

We evaluate four LLMs with different scales and access settings, by providing them the NL question and the generated non-relational views: gpt-5-mini, Qwen3.5-122B-A10B-FP8 (Qwen 122B), Qwen3.5-35B-A3B-FP8 (Qwen 35B), and Qwen3.5-9B (Qwen 9B). This covers one strong closed model and open-weight models of different sizes, testing whether failures exposed by GRADINO are model-family specific or persist across scale. All models use zero-shot Chain-of-Thought (Kojima et al., 2022), with *thinking* enabled, temperature equal to 0 and a token budget of 16384 for Qwen; details are in §A. Additional experiments with Program-of-Thoughts prompting (Chen et al., 2023) are detailed in §D.

4.2 Metrics

We report accuracy using normalized exact match (EM) at 2-decimal precision against the automatically computed gold answer, after removing superficial formatting differences, such as surrounding whitespace and thousands separators. Accuracy is the percentage of examples whose normalized prediction matches the gold answer.

4.3 Datasets

We generate GRADINO instances in two domains, environmental and finance, using gpt-5-mini as its internal LLM. Unless otherwise stated, evaluation covers parallel and sequential questions over 2, 3, and 5 tables, with parallel questions also scaled to 10 and 20 tables. We request 50 instances for




Figure 5: Performance gap between comparative and quantitative parallel questions as table count increases, averaged over domains and unit settings.




Figure 6: Per-model, per-domain accuracy drop from unit heterogeneity on parallel questions.

each combination of domain, table count, aggregation operator (sum, average, superlative), and, for parallel questions, same- or mixed-unit setting. Perturbations are applied randomly but compositionally: each larger benchmark extends the next smaller one, so 3-table instances reuse 2-table inputs, and so on. Thus, every ablation reuses the same inputs, varying only the studied property. In all experiments, rendered views are capped at 10×10 cells to control for table-size variance. The cap is a generation parameter, and larger views can be rendered by relaxing it. After discarding generations that fail the schema-format check (§3.1, §B), we obtain 2,322 parallel and 840 sequential instances; further statistics are in §E.

4.4 Discussion

(RQ1) Scaling collapses performance. As shown in Figure 4, accuracy degrades sharply as the number of tables grows, in both parallel and sequential regimes. For parallel questions, scaling from 2 to 20 tables costs gpt-5-mini 39.6 points, and Qwen 122B, 35B, and 9B 26.1, 32.8, and 38.3 points (Figure 4a). gpt-5-mini leads with few tables (95.53% at 3) but collapses to 52.61% at 20, where Qwen 122B, 35B, and 9B reach 64.90%, 56.42%, and 26.74%. Sequential accuracy follows the same trend: extending the chain from 2 to 5 tables costs 32.5 points for gpt-5-mini and 28.8, 24.6, and 10.9 points for Qwen 122B, 35B, and 9B (Figure 4b). Additionally, sequential reasoning is harder than parallel aggregation. Because it forces the model to propagate lookup keys across tables rather than combine independent evidence, sequential accuracy drops even faster per table, with the gap over parallel questions widening from 21.6

points at 2 tables to 40.1 at 5 (Figure 4c). The two regimes thus stress models differently: sequential chains are harder per table, whereas parallel questions scale to far more tables, where the collapse makes them comparably difficult.

(RQ2) Quantitative reasoning and unit heterogeneity are a major bottleneck. The scaling collapse is largely driven by precise numerical computation, as quantitative reasoning drives the collapse. The gap between the accuracy over comparative and quantitative parallel questions widens with table count for every model (Figure 5): at 20 tables it peaks at 53.6 points for gpt-5-mini, and reaches 32.0, 30.1, and 16.8 points for Qwen 122B, 35B, and 9B, so quantitative questions account for a growing share of errors. This effect is amplified by unit heterogeneity, which causes a domain-dependent collapse. As shown in Figure 6, in the environmental domain, where questions often require physical conversions (e.g., gallons to liters), it costs gpt-5-mini 50.1 points and Qwen 122B, 35B, and 9B 45.3, 46.8, and 38.9, whereas in finance, dominated by scale conversions (e.g., thousands to billions), the same perturbation is far milder (6.7, 7.3, 9.9, and 14.2 points). Robustness to numerical representation is thus domain-specific, as each unit convention induce distinct failure modes (Santacroce et al., 2026).

(RQ3) GRADINO-generated data has functional utility. Beyond evaluation, we test whether GRADINO-generated data is useful as training supervision by augmenting it with gpt-5-mini reasoning traces whose final answers match the SQL-derived gold answers, and fine-tuning Qwen2.5-14B-Instruct. We evaluate on GRI-QA (Contalbo et al., 2025), which poses parallel

Train data	2		3		5		Mean
	com	qnt	com	qnt	com	qnt	
None	18.3	9.3	17.1	9.7	9.2	0.0	10.6
GRADINO (fin)	22.3	46.4	23.3	17.7	14.7	5.8	21.7 ^{+11.1}
GRADINO (env)	35.5	50.0	35.3	18.3	25.8	5.0	28.3 ^{+17.7}
TAT-QA	33.1	46.4	28.3	19.0	13.3	4.2	24.1 ^{+13.5}
GRI-QA	73.5	44.9	53.0	15.0	32.5	3.3	37.0 ^{+26.4}
GRADINO (env) + GRI-QA	83.5	52.9	70.0	15.0	48.1	12.5	47.0 ^{+36.4}

Table 2: Accuracy of Qwen2.5-14B-Instruct fine-tuned on different training splits, evaluated on GRI-QA by number of tables and question type (comparative, quantitative), averaged over 3 seeds. Superscripts report the gain over the non-finetuned baseline; std. dev. in §D.2.

comparative and quantitative questions over 2-, 3-, and 5-table environmental settings (Table 2). Since no comparable multi-table non-relational training set exists, which is the very gap GRADINO targets, we compare against the three most relevant options: the in-domain GRI-QA train split, an out-of-domain single-table set (financial TAT-QA Zhu et al., 2021), and no fine-tuning. The TAT-QA baseline probes a common deployment scenario: a practitioner who needs a multi-table system in a target domain, but lacks in-domain multi-table data, falls back on an available single-table set from another domain. To keep the comparison fair given GRI-QA’s limited size, every regime is capped at 440 instances, a cap that is itself a symptom of the data scarcity GRADINO addresses. All regimes improve over the non-finetuned baseline (+11.1 to +36.4 on average), and two take-aways emerge. First, *in-domain multi-table supervision matters most*. Because both the domain and the reasoning structure of the training data align with the test setting, GRADINO (env) leads out-of-domain TAT-QA at every table count, with the gap widening from +3.0 at 2 tables to +6.7 at 5. TAT-QA collapses as tables accumulate (8.8 average at 5 tables), while GRADINO (env) degrades more gracefully (15.4). Second, *GRADINO complements scarce in-domain data*. As expected, GRI-QA leads GRADINO (env) overall, since it is trained on the same distribution as the test set; however, GRI-QA’s train split is skewed toward comparative questions ($\sim 3\times$ more comparative than quantitative), so the GRI-QA-tuned model stays weak on quantitative reasoning compared to GRADINO (env). Moreover, GRADINO can cheaply supply the multi-table examples GRI-QA’s split lacks: com-

Topology	Quantitative		Comparative	
	Avg	Sum	Superlative	Mean
Parallel	100.0	100.0	97.5	99.2
Sequential	95.0	87.5	95.0	92.5
Mean	97.5	93.8	96.3	95.9

Table 3: Manually verified correctness (%) of GRADINO-generated instances.

binning GRI-QA + GRADINO (env) yields the best model overall (+36.4) and improves performance at every table count and question type over GRI-QA alone. This confirms that GRADINO augments rather than merely substitutes for real data, providing the scalable, multi-table supervision that existing benchmarks lack.

(RQ4) Question generation is reliable. Question generation is the only GRADINO stage that can introduce errors, since answers and evidence are guaranteed correct by construction. To estimate its reliability, we manually inspected 240 instances in the environmental and finance domain, 10 per aggregation operator (sum, average, superlative) across parallel (20 and 10 tables, unit-converted) and sequential (5 and 3 tables) regimes. Since we target the most error-prone settings, with many tables and unit conversions, reliability is likely higher in easier ones. *Reliability is high*, as parallel instances average 99.2%, while sequential ones 92.5% (Table 3). The gap is expected, because sequential verbalization must coordinate constraints across heterogeneous tables rather than reuse a shared predicate. All errors trace to a dropped or hallucinated WHERE constraint, likely fixable with a stronger backbone than gpt-5-mini. Moreover, *this reliability is competitive with manually curated benchmarks*: even in the hardest settings, GRADINO’s measured error rate stays below those reported for hand-built tabular datasets (Park et al., 2026), where error rates can reach up to 30%.

5 Conclusion

We introduced GRADINO, a framework for generating numerical QA benchmarks over multiple non-relational tables, giving controlled access to a setting that existing resources cannot vary. Using it, we identify multi-table non-relational numerical reasoning as a persistent blind spot for current LLMs, which degrade sharply as tables accumulate and with domain-dependent unit conventions.

Moreover, we show that GRADINO can generate useful data for LLM fine-tuning.

6 Limitations

Synthetic data realism. Our experiments on perturbed real tables validate the observed degradation on generated tables. However, GRADINO is designed as a controlled diagnostic generator, not as a replacement for manually curated benchmarks built from naturally occurring documents. Although the generator is conditioned on domain descriptions, unit inventories, and optional examples, the resulting tables may not capture all distributional, visual, and semantic regularities of real spreadsheets or reports, not even with GRADINO applied in a semi-automatic manner (§B.4). This is a deliberate trade-off: from-scratch generation gives us control over table count, reasoning topology, perturbations, and units, but it does not guarantee full realism.

Restricted numerical structure. This study focuses on numerical QA, where each latent table contains a distinguished numerical value attribute. Although GRADINO can also generate non-numerical questions, we restrict our analysis to numerical questions because parallel scenarios, which require evidence retrieval from multiple tables followed by an aggregation operation, are especially common with numerical values. Future work could extend GRADINO to non-numerical settings, including questions that require set operations, categorical comparisons, or logical aggregation over textual evidence.

Question verbalization. Gold answers are computed by executing latent SQL programs, and non-relational views are generated through answer-preserving transformations. However, the NL question is produced by an LLM and is therefore not guaranteed correct by construction. We mitigate this issue through verifier-based repair and manual inspection. Table 3 shows high reliability overall, but also that sequential questions are more error-prone, mainly because verbalization may drop latent selection constraints. Thus, GRADINO provides executable guarantees for answers and evidence preservation, while question fidelity remains only an empirically validated component.

Generation failures and filtering. Some generations fail schema-format checks. We discard invalid generations, which improves benchmark quality but increases the cost of data generation. Future work should focus on reducing the number of

discarded generations to maximize QA generation throughput.

Perturbation coverage. We consider a broad set of structural, cell-level, layout, and numerical perturbations. Nevertheless, real non-relational tables contain additional phenomena that we do not fully model, such as visual alignment cues and chart-table mixtures. Although GRADINO could be extended to cover those perturbations, the current benchmark does not stress the full range of document-table understanding. We leave this as future work.

7 Ethical considerations

Data and privacy. GRADINO generates its evaluation data entirely from scratch: latent relational sources, tables, and questions are synthesized by LLMs from a domain specification and contain only fictional entities and randomly sampled numerical values. The generated instances therefore include no personally identifying information and no real individuals, and we collect no scraped or user-contributed content. The real tables used in our validation experiments (§C) come from public databases used exclusively for evaluation (Northwind, BikeStores, SQL Bootcamp) and from existing research benchmarks (Spider, BIRD, BEAVER, TAT-QA, GRI-QA) built from public corporate and web sources. Where such sources mention named public figures, this information is already public and is neither augmented nor inferred by GRADINO.

Use of existing artifacts. We use all datasets and models in a manner consistent with their intended research use, and we report the license of each artifact in §G. Our use is non-commercial and research-only, and any derivatives we produce remain within a research context. The crawled sample databases are used solely for testing, without redistribution (§C).

Risks and broader impact. GRADINO is a diagnostic tool for measuring the robustness of Table QA systems, and we see limited potential for direct misuse. The principal risk is interpretive: because the benchmark is synthetic and controlled, strong performance on GRADINO should not be read as a guarantee of real-world reliability, nor should low scores be taken as evidence that a model is unusable in practice. We therefore position GRADINO as a complement to, rather than a replacement for,

manually curated benchmarks built from naturally occurring documents (§6). Used appropriately, surfacing the multi-table, non-relational, and unit-heterogeneity failure modes, we believe GRADINO can support the development of more reliable analytical interfaces over tabular data, including in financial and environmental reporting settings where silent numerical errors carry real consequences.

Human verification and AI assistants. The correctness study in Table 3 was carried out by the authors and involves no external annotators or human subjects, raising no participant-compensation or consent concerns. Our use of AI assistants for writing is described in §8, and the role of LLMs within the generation pipeline is described in §3.

8 Use of AI assistants

When writing this paper, we used ChatGPT and Claude to improve the flow of writing and the vocabulary of the initial drafts we manually wrote. Each suggestion has been manually validated by the authors.

References

- QUDT: Quantities, units, dimensions and types ontologies. <https://www.qudt.org/>. Accessed: 2026-05-09.
- Mohammad Shahmeer Ahmad, Zan Ahmad Naeem, Michaël Aupetit, Ahmed K. Elmagarmid, Mohamed Y. Eltabakh, Xiaosong Ma, Mourad Ouzzani, and Chaoyi Ruan. 2025. [HCT-QA: A benchmark for question answering on human-centric tables](#). *CoRR*, abs/2504.20047.
- Denver Baumgartner and Tomasz Kornuta. 2024. [SynQL: Synthetic data generation for in-domain, low-resource text-to-SQL parsing](#). In *NeurIPS 2024 Third Table Representation Learning Workshop*.
- Shuaichen Chang, Jun Wang, Mingwen Dong, Lin Pan, Henghui Zhu, Alexander Hanbo Li, Wuwei Lan, Sheng Zhang, Jiarong Jiang, Joseph Lilien, Steve Ash, William Yang Wang, Zhiguo Wang, Vittorio Castelli, Patrick Ng, and Bing Xiang. 2023. [Dr.spider: A diagnostic evaluation benchmark towards text-to-SQL robustness](#). In *The Eleventh International Conference on Learning Representations*.
- Peter Baile Chen, Fabian Wenz, Yi Zhang, Moe Kayali, Nesime Tatbul, Michael J. Cafarella, Çagatay Demiralp, and Michael Stonebraker. 2024. [BEAVER: an enterprise benchmark for text-to-sql](#). *CoRR*, abs/2409.02038.
- Wenhu Chen, Ming-Wei Chang, Eva Schlinger, William Yang Wang, and William W. Cohen. 2021a.

[Open question answering over tables and text](#). In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. OpenReview.net.

Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. 2023. [Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks](#). *Transactions on Machine Learning Research*.

Wenhu Chen, Hongmin Wang, Jianshu Chen, Yunkai Zhang, Hong Wang, Shiyang Li, Xiyu Zhou, and William Yang Wang. 2020a. [Tabfact: A large-scale dataset for table-based fact verification](#). In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.

Wenhu Chen, Hanwen Zha, Zhiyu Chen, Wenhan Xiong, Hong Wang, and William Yang Wang. 2020b. [Hybridqa: A dataset of multi-hop question answering over tabular and textual data](#). In *Findings of the Association for Computational Linguistics: EMNLP 2020, Online Event, 16-20 November 2020*, Findings of ACL, pages 1026–1036. Association for Computational Linguistics.

Zhiyu Chen, Wenhu Chen, Charese Smiley, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Kenneth Huang, Bryan R. Routledge, and William Yang Wang. 2021b. [Finqa: A dataset of numerical reasoning over financial data](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 7-11 November, 2021*, pages 3697–3711. Association for Computational Linguistics.

Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma, Sameena Shah, and William Yang Wang. 2022. [Convfinqa: Exploring the chain of numerical reasoning in conversational finance question answering](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 6279–6292. Association for Computational Linguistics.

Zhoujun Cheng, Haoyu Dong, Zhiruo Wang, Ran Jia, Jiaqi Guo, Yan Gao, Shi Han, Jian-Guang Lou, and Dongmei Zhang. 2022. [Hitab: A hierarchical table dataset for question answering and natural language generation](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 1094–1110. Association for Computational Linguistics.

Xu Chu, Ihab F. Ilyas, and Paolo Papotti. 2013. [Discovering denial constraints](#). *Proc. VLDB Endow.*, 6(13):1498–1509.

Michele Luca Contalbo, Sara Pederzoli, Francesco Del Buono, Venturilli Valeria, Francesco Guerra, and

- Matteo Paganelli. 2025. [GRI-QA: a comprehensive benchmark for table question answering over environmental data](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 15764–15779, Vienna, Austria. Association for Computational Linguistics.
- Leonardo De Moura and Nikolaj Bjørner. 2008. Z3: an efficient smt solver. In *Proceedings of the Theory and Practice of Software, 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems, TACAS’08/ETAPS’08*, page 337–340, Berlin, Heidelberg. Springer-Verlag.
- Yang Deng, Wenqiang Lei, Wenxuan Zhang, Wai Lam, and Tat-Seng Chua. 2022. [PACIFIC: towards proactive conversational question answering over tabular and textual data in finance](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 6970–6984. Association for Computational Linguistics.
- Shaoming Duan, Youxuan Wu, Chuanyi Liu, Yuhao Zhang, Zirui Wang, Peiyi Han, Shengyuan Yu, Liang Yan, and Yingwei Liang. 2025. [DSQG-syn: Synthesizing high-quality data for text-to-SQL parsing by domain specific question generation](#). In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 2971–2989, Albuquerque, New Mexico. Association for Computational Linguistics.
- Liancheng Fang, Aiwei Liu, Hengrui Zhang, Henry Peng Zou, Weizhi Zhang, and Philip S. Yu. 2025. [TABGEN-ICL: Residual-aware in-context example selection for tabular data generation](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 20027–20041, Vienna, Austria. Association for Computational Linguistics.
- Jonathan Fürst, Catherine Kosten, Farhad Nooralahzadeh, Yi Zhang, and Kurt Stockinger. 2025. [Evaluating the data model robustness of text-to-sql systems based on real user queries](#). In *Proceedings 28th International Conference on Extending Database Technology, EDBT 2025, Barcelona, Spain, March 25-28, 2025*, pages 158–170. OpenProceedings.org.
- Yujian Gan, Xinyun Chen, Qiuping Huang, Matthew Purver, John R. Woodward, Jinxia Xie, and Pengsheng Huang. 2021. [Towards robustness of text-to-SQL models against synonym substitution](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 2505–2515, Online. Association for Computational Linguistics.
- Akash Ghosh, Venkata Sahith Bathini, Niloy Ganguly, Pawan Goyal, and Mayank Singh. 2024. [How robust are the QA models for hybrid scientific tabular data? a study using customized dataset](#). In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 8258–8264, Torino, Italia. ELRA and ICCL.
- Yu Guo, Dong Jin, Shenghao Ye, Shuangwu Chen, Jian Yang, and Xiaobin Tan. 2025. [SQLForge: Synthesizing reliable and diverse data to enhance text-to-SQL reasoning in LLMs](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 8441–8452, Vienna, Austria. Association for Computational Linguistics.
- Jonathan Herzig, Thomas Müller, Syrine Krichene, and Julian Martin Eisenschlos. 2021. [Open domain question answering over tables via dense retrieval](#). In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2021, Online, June 6-11, 2021*, pages 512–519. Association for Computational Linguistics.
- Frederik Hoppe, Astrid Franz, Lars Kleinemeier, and Udo Göbel. 2025. [Generating synthetic relational tabular data via structural causal models](#). In *Proceedings of the 4th Table Representation Learning Workshop*, pages 13–18, Vienna, Austria. Association for Computational Linguistics.
- Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [LoRA: Low-rank adaptation of large language models](#). In *International Conference on Learning Representations*.
- Zhengbao Jiang, Yi Mao, Pengcheng He, Graham Neubig, and Weizhu Chen. 2022. [OmniTab: Pretraining with natural and synthetic data for few-shot table-based question answering](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 932–942, Seattle, United States. Association for Computational Linguistics.
- Yannis Katsis, Saneem Chemmengath, Vishwajeet Kumar, Samarth Bharadwaj, Mustafa Canim, Michael Glass, Alfio Gliozzo, Feifei Pan, Jaydeep Sen, Karthik Sankaranarayanan, and Soumen Chakrabarti. 2022. [AIT-QA: Question answering dataset over complex tables in the airline industry](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies: Industry Track*, pages 305–314, Hybrid: Seattle, Washington + Online. Association for Computational Linguistics.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. [Large language models are zero-shot reasoners](#). In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*.
- Jinyang Li, Binyuan Hui, GE QU, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying

894	Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guo-	generations. In <i>The Thirty-eighth Annual Conference</i>	951
895	liang Li, Kevin Chang, Fei Huang, Reynold Cheng,	on Neural Information Processing Systems.	952
896	and Yongbin Li. 2023. Can LLM already serve as		
897	a database interface? a BIG bench for large-scale		
898	database grounded text-to-SQLs . In <i>Thirty-seventh</i>		
899	<i>Conference on Neural Information Processing Sys-</i>		
900	<i>tems Datasets and Benchmarks Track</i> .		
901	Qian Liu, Bei Chen, Jiaqi Guo, Morteza Ziyadi, Zeqi	Simone Papicchio, Paolo Papotti, and Luca Cagliero.	953
902	Lin, Weizhu Chen, and Jian-Guang Lou. 2022.	2023. QATCH: benchmarking sql-centric tasks with	954
903	TAPEX: table pre-training via learning a neural SQL	table representation learning models on your data . In	955
904	executor . In <i>The Tenth International Conference on</i>	<i>Advances in Neural Information Processing Systems</i>	956
905	<i>Learning Representations, ICLR 2022, Virtual Event,</i>	<i>36: Annual Conference on Neural Information Pro-</i>	957
906	<i>April 25-29, 2022</i> . OpenReview.net.	<i>cessing Systems 2023, NeurIPS 2023, New Orleans,</i>	958
		<i>LA, USA, December 10 - 16, 2023</i> .	959
907	Tennison Liu, Zhaozhi Qian, Jeroen Berrevoets, and	Sungho Park, Jueun Kim, and Wook-Shin Han. 2026.	960
908	Mihaela van der Schaar. 2023. GOGGLE: Genera-	SPARTA: Scalable and principled benchmark of tree-	961
909	tive modelling for tabular data by learning relational	structured multi-hop QA over text and tables . In	962
910	structure . In <i>The Eleventh International Conference</i>	<i>The Fourteenth International Conference on Learn-</i>	963
911	<i>on Learning Representations</i> .	<i>ing Representations</i> .	964
912	Tianyang Liu, Fei Wang, and Muhao Chen. 2024. Re-	Panupong Pasupat and Percy Liang. 2015. Compo-	965
913	thinking tabular data understanding with large lan-	sitional semantic parsing on semi-structured tables .	966
914	guage models . In <i>Proceedings of the 2024 Confer-</i>	In <i>Proceedings of the 53rd Annual Meeting of the</i>	967
915	<i>ence of the North American Chapter of the Associ-</i>	<i>Association for Computational Linguistics and the</i>	968
916	<i>ation for Computational Linguistics: Human Lan-</i>	<i>7th International Joint Conference on Natural Lan-</i>	969
917	<i>guage Technologies (Volume 1: Long Papers)</i> , pages	<i>guage Processing of the Asian Federation of Natural</i>	970
918	450–482, Mexico City, Mexico. Association for Com-	<i>Language Processing, ACL 2015, July 26-31, 2015,</i>	971
919	putational Linguistics.	<i>Beijing, China, Volume 1: Long Papers</i> , pages 1470–	972
		1480. The Association for Computer Linguistics.	973
920	Pan Lu, Liang Qiu, Kai-Wei Chang, Ying Nian Wu,	Xinyu Pi, Bing Wang, Yan Gao, Jiaqi Guo, Zhoujun	974
921	Song-Chun Zhu, Tanmay Rajpurohit, Peter Clark,	Li, and Jian-Guang Lou. 2022. Towards robustness	975
922	and Ashwin Kalyan. 2023. Dynamic prompt learning	of text-to-SQL models against natural and realistic	976
923	via policy gradient for semi-structured mathematical	adversarial table perturbation . In <i>Proceedings of the</i>	977
924	reasoning . In <i>The Eleventh International Conference</i>	<i>60th Annual Meeting of the Association for Compu-</i>	978
925	<i>on Learning Representations</i> .	<i>tational Linguistics (Volume 1: Long Papers)</i> , pages	979
		2007–2022, Dublin, Ireland. Association for Compu-	980
		tational Linguistics.	981
926	Linyong Nan, Chiachun Hsieh, Ziming Mao, Xi Victo-	Zipeng Qiu, You Peng, Guangxin He, Binhang Yuan,	982
927	ria Lin, Neha Verma, Rui Zhang, Wojciech Kryscin-	and Chen Wang. 2024. Tqa-bench: Evaluating llms	983
928	ski, Hailey Schoelkopf, Riley Kong, Xiangru Tang,	for multi-table question answering with scalable con-	984
929	Mutethia Mutuma, Ben Rosand, Isabel Trindade,	text and symbolic extension . <i>CoRR</i> , abs/2411.19504.	985
930	Renusree Bandaru, Jacob Cunningham, Caiming		
931	Xiong, and Dragomir R. Radev. 2022. Fetaqa: Free-	Joeri Rammelaere and Floris Geerts. 2019. Revisiting	986
932	form table question answering . <i>Trans. Assoc. Com-</i>	conditional functional dependency discovery: Split-	987
933	<i>put. Linguistics</i> , 10:35–49.	ting the “c” from the “fd”. In <i>Machine Learning and</i>	988
		<i>Knowledge Discovery in Databases</i> , pages 552–568,	989
		Cham. Springer International Publishing.	990
934	Jio Oh, Soyeon Kim, Junseok Seo, Jindong Wang,	Federico Ranaldi, Elena Sofia Ruzzetti, Dario Ono-	991
935	Ruochen Xu, Xing Xie, and Steven Whang. 2024.	rati, Leonardo Ranaldi, Cristina Giannone, Andrea	992
936	Erbench: An entity-relationship based automatically	Favalli, Raniero Romagnoli, and Fabio Massimo Zan-	993
937	verifiable hallucination benchmark for large language	zotto. 2024. Investigating the impact of data contami-	994
938	models . In <i>Advances in Neural Information Process-</i>	nation of large language models in text-to-SQL trans-	995
939	<i>ing Systems 38: Annual Conference on Neural In-</i>	lation . In <i>Findings of the Association for Computa-</i>	996
940	<i>formation Processing Systems 2024, NeurIPS 2024,</i>	<i>tational Linguistics: ACL 2024</i> , pages 13909–13920,	997
941	<i>Vancouver, BC, Canada, December 10 - 15, 2024</i> .	Bangkok, Thailand. Association for Computational	998
		Linguistics.	999
942	Vaishali Pal, Andrew Yates, Evangelos Kanoulas, and	Varshini Reddy, Rik Koncel-Kedziorski, Viet Dac Lai,	1000
943	Maarten de Rijke. 2023. MultiTabQA: Generating	Michael Krumdick, Charles Lovering, and Chris Tan-	1001
944	tabular answers for multi-table question answering .	ner. 2024. DocFinQA: A long-context financial rea-	1002
945	In <i>Proceedings of the 61st Annual Meeting of the</i>	soning dataset . In <i>Proceedings of the 62nd Annual</i>	1003
946	<i>Association for Computational Linguistics (Volume</i>	<i>Meeting of the Association for Computational Lin-</i>	1004
947	<i>1: Long Papers)</i> , pages 6322–6334, Toronto, Canada.	<i>guistics (Volume 2: Short Papers)</i> , pages 445–458,	1005
948	Association for Computational Linguistics.	Bangkok, Thailand. Association for Computational	1006
949	Arjun Panickssery, Samuel R. Bowman, and Shi Feng.	Linguistics.	1007
950	2024. LLM evaluators recognize and favor their own		

1008	Marta Santacroce, Michele Luca Contalbo, Sara Ped-	diverse reasoning types. In <i>Findings of the Asso-</i>	1065
1009	erzoli, Riccardo Benassi, Venturelli Valeria, Matteo	ciation for Computational Linguistics: <i>ACL 2025</i> ,	1066
1010	Paganelli, and Francesco Guerra. 2026. CLARIESG:	pages 3859–3881, Vienna, Austria. Association for	1067
1011	An end-to-end system for ESG analysis over com-	Computational Linguistics.	1068
1012	plex tables in corporate reports. In <i>Proceedings of</i>		
1013	<i>the 19th Conference of the European Chapter of the</i>	Yilun Zhao, Yunxiang Li, Chenying Li, and Rui Zhang.	1069
1014	<i>Association for Computational Linguistics (Volume</i>	2022. MultiHiertt: Numerical reasoning over multi	1070
1015	<i>3: System Demonstrations)</i> , pages 86–100, Rabat,	hierarchical tabular and textual data. In <i>Proceedings</i>	1071
1016	Marocco. Association for Computational Linguistics.	<i>of the 60th Annual Meeting of the Association for</i>	1072
1017	Aivin V. Solatorio and Olivier Dupriez. 2023. Realtab-	<i>Computational Linguistics (Volume 1: Long Papers)</i> ,	1073
1018	former: Generating realistic relational and tabular	pages 6588–6600, Dublin, Ireland. Association for	1074
1019	data using transformers. <i>CoRR</i> , abs/2302.02041.	Computational Linguistics.	1075
1020	Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa	Victor Zhong, Caiming Xiong, and Richard Socher.	1076
1021	Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh	2017. Seq2sql: Generating structured queries	1077
1022	Hajishirzi. 2023. Self-instruct: Aligning language	from natural language using reinforcement learning.	1078
1023	models with self-generated instructions. In <i>Proceed-</i>	<i>CoRR</i> , abs/1709.00103.	1079
1024	<i>ings of the 61st Annual Meeting of the Association for</i>		
1025	<i>Computational Linguistics (Volume 1: Long Papers)</i> ,	Wei Zhou, Mohsen Mesgar, Heike Adel, and Annemarie	1080
1026	pages 13484–13508, Toronto, Canada. Association	Friedrich. 2024. FREB-TQA: A fine-grained robust-	1081
1027	for Computational Linguistics.	ness evaluation benchmark for table question answer-	1082
1028	Jian Wu, Linyi Yang, Dongyuan Li, Yuliang Ji, Manabu	ing. In <i>Proceedings of the 2024 Conference of the</i>	1083
1029	Okumura, and Yue Zhang. 2025a. MMQA: Evaluat-	<i>North American Chapter of the Association for Com-</i>	1084
1030	ing LLMs with multi-table multi-hop complex ques-	<i>putational Linguistics: Human Language Technolo-</i>	1085
1031	tions. In <i>The Thirteenth International Conference on</i>	<i>gies (Volume 1: Long Papers)</i> , NAACL 2024, Mexico	1086
1032	<i>Learning Representations.</i>	<i>City, Mexico, June 16-21, 2024</i> , pages 2479–2497.	1087
1033	Pengzuo Wu, Yuhang Yang, Guangcheng Zhu, Chao	Association for Computational Linguistics.	1088
1034	Ye, Hong Gu, Xu Lu, Ruixuan Xiao, Bowen Bao,		
1035	Yijing He, Liangyu Zha, Wentao Ye, Junbo Zhao,	Fengbin Zhu, Wenqiang Lei, Youcheng Huang, Chao	1089
1036	and Haobo Wang. 2025b. Realhitbench: A com-	Wang, Shuo Zhang, Jiancheng Lv, Fuli Feng, and	1090
1037	prehensive realistic hierarchical table benchmark for	Tat-Seng Chua. 2021. TAT-QA: A question answer-	1091
1038	evaluating llm-based table analysis. In <i>Findings of</i>	ing benchmark on a hybrid of tabular and textual	1092
1039	<i>the Association for Computational Linguistics, ACL</i>	content in finance. In <i>Proceedings of the 59th An-</i>	1093
1040	2025, Vienna, Austria, July 27 - August 1, 2025, Find-	<i>annual Meeting of the Association for Computational</i>	1094
1041	ings of ACL, pages 7105–7137. Association for Com-	<i>Linguistics and the 11th International Joint Confer-</i>	1095
1042	putational Linguistics.	<i>ence on Natural Language Processing, ACL/IJCNLP</i>	1096
1043	XBRL International. XBRL units registry 1.0. https://specifications.xbrl.org/work-product-i-	2021, (Volume 1: Long Papers), Virtual Event, Au-	1097
1044	ndex-registries-units-registry-1.0.html .	gust 1-6, 2021, pages 3277–3287. Association for	1098
1045	Accessed: 2026-05-09.	Computational Linguistics.	1099
1046			
1047	Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga,	A Appendix Overview	1100
1048	Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingn-	This appendix provides additional methodological	1101
1049	ing Yao, Shanelle Roman, Zilin Zhang, and Dragomir	details and extended results complementing the	1102
1050	Radev. 2018. Spider: A large-scale human-labeled	main paper.	1103
1051	dataset for complex and cross-domain semantic pars-	Appendix B details the lightweight semantic con-	1104
1052	ing and text-to-SQL task. In <i>Proceedings of the 2018</i>	straints referenced in §3.1, which GRADINO can	1105
1053	<i>Conference on Empirical Methods in Natural Lan-</i>	optionally enforce while instantiating the latent re-	1106
1054	<i>guage Processing</i> , pages 3911–3921, Brussels, Bel-	lational source to push it toward more plausible	1107
1055	gium. Association for Computational Linguistics.	configurations.	1108
1056	Tianshu Zhang, Kun Qian, Siddhartha Sahai, Yuan	Appendix C reports the main tests cited in §1 that	1109
1057	Tian, Shaddy Garg, Huan Sun, and Yunyao Li.	motivate GRADINO. In particular, it shows that per-	1110
1058	2025a. Evoschema: Towards text-to-sql robust-	turbed parallel and sequential questions are difficult	1111
1059	ness against schema evolution. <i>Proc. VLDB Endow.</i> ,	to synthesize from existing relational tables, and	1112
1060	18(10):3655–3668.	that non-relationalization markedly degrades LLM	1113
1061	Xuanliang Zhang, Dingzirui Wang, Baoxin Wang,	performance in multi-table settings.	1114
1062	Longxu Dou, Xinyuan Lu, Keyan Xu, Dayong Wu,	Appendix D reports the detailed experimental re-	1115
1063	and Qingfu Zhu. 2025b. SCITAT: A question answer-	sults underlying §4.4, together with the finetuning	1116
1064	ing benchmark for scientific tables and text covering	configuration. It provides per-setting performance	1117
		of all evaluated baselines on parallel and sequential	1118

questions over both relational and non-relational tables, as well as the training setup and data composition used in the finetuning experiments.

Appendix E reports statistics on the GRADINO-generated data used throughout the experiments, including the number of retained instances per experiment, the share of candidates discarded during generation, and how instances are balanced across table counts, aggregation operators, unit settings, and relational versus non-relational views.

Appendix F illustrates the prompts used at each stage of GRADINO, covering latent relational source generation, semantic constraint generation, natural-language question generation and verification for both parallel and sequential regimes, and the inference prompt given to the evaluated models.

Appendix G details the licenses of all artifacts used in this work, all of which permit research use.

B Enforcing semantic constraints

This appendix describes the lightweight semantic regularities referenced in §3.1, which GRADINO can optionally enforce while instantiating the latent relational source. The goal is to push the densely sampled source toward more plausible configurations. These constraints are not part of the supervision signal and are not evaluated quantitatively; gold answers are always computed by executing the SQL supervision over the instantiated tables. The process has three steps: an LLM *defines* constraints over the schema (§B.1), the constraints are *parsed* into a solver representation (§B.2), and feasible values are *sampled* subject to them (§B.3). §B.4 further demonstrates a usage mode in which the user semi-automatically defines and validates the schema and constraints prior to QA generation with GRADINO. In this setting, constraint plausibility is determined by the user rather than left to the LLM.

B.1 Rule Definition

We adopt two constraint families. **Intra-row constraints** (Conditional Functional Dependencies Rammelaere and Geerts, 2019) bind semantically related values within a tuple: whenever a row matches a selector, one of its attributes must satisfy an (in)equality (e.g., "*if the city is Lisbon, the country must be Portugal*"). **Inter-row constraints** (Denial Constraints Chu et al., 2013) express order or arithmetic relations on the numerical value column v across rows matched on a selector (e.g.,

"*net income must always be lower than the gross income*"). Given the schema metadata from §3.1, an LLM is prompted to infer realistic rules for both constraint families using its internal domain knowledge. Rules use a restricted, Python-evaluable syntax: a row is selected by fixing categorical attributes, e.g. `(TypeOfMoney == "Gross")`, with all unspecified attributes implicitly held equal across matched rows. Intra-row rules take the conditional form `if (selector) then (predicate)`, e.g. `if (CompanyTier == "Small") then (Amount <= 5000000)`, while inter-row rules compare the value column of two selected rows through the attribute-as-property syntax `(selector).v`, e.g. `(TypeOfMoney == "Gross").Amount > (TypeOfMoney == "Net").Amount`. More information on the rules format and allowed operations can be found inside the prompt shown in Figure 10.

B.2 Parsing into a Solver Representation

Each rule is translated into a constraint over the Z3 SMT solver (De Moura and Bjørner, 2008), using one solver variable per (attribute, row) pair. Categorical attributes are encoded as enumerated sorts over their candidate values and numerical attributes as integer or real variables. Intra-row rules become implications from the selector to the predicate, applied to the rows the selector matches; inter-row rules are grounded by substituting the value variables of the two matched rows into the comparison. The declared range of v , defined during schema generation, is added as an explicit bound.

B.3 Randomized Value Sampling

Enforcing the constraints must not collapse the table to a single canonical solution, so we sample values within their feasible region rather than returning an arbitrary satisfying model. For intra-row constraints, for each variable we use Z3 to compute its feasible interval $[\min, \max]$ given all constraints and previously fixed variables. Then, we draw a value uniformly from that interval, assign it to the variable, and then add the novel constraint inside the optimizer’s constraints. For inter-row constraints, since each rule may contain multiple numerical variables belonging to different rows, optimizing all the variables may become computationally expensive. For this reason, for each variable we instead obtain two feasible values with satisfiability queries and treat them as $[\min, \max]$ values.

Because each variable is bounded conditionally

		2			3			5			10			20		
		com	quant	avg	com	quant	avg	com	quant	avg	com	quant	avg	com	quant	avg
environmental	same unit															
	gpt-5-mini	94.74	89.47	92.11	92.11	93.43	92.77	95.45	89.19	92.32	87.10	60.75	73.93	87.88	21.21	54.55
	Qwen3.5-122B	89.47	96.06	92.76	97.37	89.47	93.42	86.47	85.14	85.80	80.56	70.84	75.70	84.85	57.58	71.22
	Qwen3.5-35B	86.84	86.84	86.84	84.21	93.43	88.82	89.19	82.44	85.81	88.89	58.34	73.61	78.79	36.36	57.58
	Qwen3.5-9B	65.79	67.11	66.45	65.79	67.11	66.45	59.46	54.05	56.76	52.78	33.33	43.06	42.42	10.61	26.51
	diff unit															
	gpt-5-mini	63.16	47.37	55.27	42.11	34.21	38.16	31.82	16.22	24.02	45.16	4.62	24.89	24.24	1.52	12.88
	Qwen3.5-122B	60.53	51.32	55.92	50.00	42.11	46.06	56.76	25.68	41.22	38.89	13.89	26.39	36.36	9.09	22.73
finance	same unit															
	gpt-5-mini	90.48	94.05	92.27	100.00	97.62	98.81	88.10	90.48	89.29	81.08	70.73	75.91	81.58	19.74	50.66
	Qwen3.5-122B	88.10	90.48	89.29	88.10	83.34	85.72	88.10	77.38	82.74	87.80	59.76	73.78	73.68	43.47	58.58
	Qwen3.5-35B	90.48	92.86	91.67	90.48	80.95	85.72	78.57	80.95	79.76	92.68	54.88	73.78	71.05	39.48	55.26
	Qwen3.5-9B	57.14	70.24	63.69	50.00	67.86	58.93	64.29	50.00	57.15	53.66	34.15	43.91	36.84	17.11	26.97
	diff unit															
	gpt-5-mini	80.95	85.72	83.33	95.24	89.29	92.26	92.86	82.14	87.50	72.97	58.54	65.75	76.32	13.16	44.74
	Qwen3.5-122B	85.71	80.95	83.33	90.48	80.95	85.72	80.95	63.10	72.02	75.61	45.13	60.37	73.68	30.26	51.97
	same unit															
	Qwen3.5-35B	80.95	77.38	79.17	90.48	78.57	84.53	71.43	66.67	69.05	73.17	47.56	60.37	55.26	31.58	43.42
	diff unit															
	Qwen3.5-9B	47.62	48.81	48.22	59.52	44.05	51.78	52.38	38.10	45.24	31.71	15.85	23.78	18.42	2.63	10.53

Table 4: Performance on parallel questions across domains, unit settings, models, and numbers of tables. For each number of tables, results are reported for comparative questions, quantitative questions, and their average.

on earlier choices, the final assignment is jointly consistent while remaining randomized. When the solution space is not continuous (e.g., constraints involving modulo), we do not apply randomized sampling. Rows whose constraints are unsatisfiable are dropped, and the sampled values are written back rounded to the precision declared during schema generation.

B.4 Enforcing constraints semi-automatically

The constraints described in §B.1–§B.3 are inferred automatically by an LLM and are therefore best-effort, since the model may miss domain-relevant regularities or propose relationships that a practitioner would consider implausible. This is acceptable in our default setting, whose goal is controlled diagnostic generation rather than faithful reproduction of real sources. When higher plausibility is desired, GRADINO exposes the same constraint machinery as a user-driven mode, following the semi-automatic supervision strategy of Ahmad et al. (2025). In this mode the LLM only proposes a candidate schema (§B.1) and a candidate constraint set in the format of Figures 9 and 10. A domain expert then inspects, edits, removes, or adds rules directly in this format, which is human-readable by design: row selectors are equalities over categorical attributes, and predicates use a restricted, Python-evaluable syntax. Authoring is incremen-

tal, as the user iterates on the proposed rules rather than writing them from scratch, and any edited rule set is passed through the same parser (§B.2) and solver-backed sampler (§B.3) as the automatic pipeline. This places domain knowledge directly under user control. GRADINO builds the interface on two well-studied constraint families, namely conditional functional dependencies for intra-row rules (Rammelaere and Geerts, 2019) and denial constraints for inter-row rules (Chu et al., 2013), chosen for their expressiveness over the categorical-and-single-numeric structure we target. The interface is extensible to further families without altering the rest of the pipeline, and the Z3 encoding of §B.2 guarantees that the instantiated latent source satisfies every rule the user supplies. Constraint coverage and plausibility thus rest with the user, who is the appropriate judge of domain validity.

Semi-automatic enforcement changes only the plausibility of the latent source. Constraints are never part of the supervision signal, and gold answers are always obtained by executing the SQL programs over the instantiated tables (§3.2), so neither the choice nor the coverage of constraints can affect answer correctness or evidence preservation. The semi-automatic mode therefore lets users trade generation effort for realism along a continuous range, from fully automatic from-scratch generation to expert-validated schemas and constraints,

		2			3			5		
		com	quant	avg	com	quant	avg	com	quant	avg
environ.	gpt-5-mini	98.00	80.00	89.00	69.39	61.23	65.31	56.25	47.92	52.08
	Qwen3.5-122B	78.00	65.00	71.50	61.22	41.84	51.53	56.25	23.96	40.11
	Qwen3.5-35B	84.00	72.00	78.00	65.31	46.94	56.13	60.42	28.13	44.27
	Qwen3.5-9B	54.00	34.00	44.00	34.69	24.49	29.59	33.33	18.75	26.04
finance	gpt-5-mini	80.43	65.22	72.83	68.89	45.56	57.22	45.24	44.05	44.65
	Qwen3.5-122B	65.22	59.78	62.50	64.44	27.78	46.11	42.86	29.76	36.31
	Qwen3.5-35B	69.57	44.57	57.07	46.67	25.56	36.11	50.00	33.33	41.67
	Qwen3.5-9B	34.78	20.65	27.72	33.33	12.22	22.78	33.33	14.29	23.81

Table 5: Performance across domains, models, and numbers of tables. For each number of tables, results are reported for comparative questions, quantitative questions, and their average.




Figure 7: Feasibility of constructing dense pivoted non-relational views from existing relational tables. Dense hierarchical pivots are rare: at density 1.0, only 7 tables support at least one header level, and only 2 support at least two header levels.




Figure 8: Performance gap of parallel questions over relational and non-relational table views. In every setting, non-relationalization degrades performance, with the drop being more significant on the multi-table setting.

while preserving GRADINO’s executable guarantees throughout.

C Generating non-relational QA instances from real relational tables

This appendix elaborates on the two main motivations underlying GRADINO: (i) the intrinsic difficulty of synthesizing non-relational QA instances from existing relational tables, and (ii) the marked drop in LLM performance when answering questions over such non-relational tables synthesized from real ones.

Perturbed parallel and sequential questions are difficult to synthesize from real tables. Not all perturbations and numerical question types can be obtained by transforming real relational tables. To test this, we extracted databases from Spider, BIRD, and BEAVER, obtaining 2,096 relational

tables, and retained only the 148 tables with at least one floating-point attribute, and satisfying tabular size constraints (more than 2 columns and 9 rows). For each table, we attempted to construct a pivoted view by selecting the attributes that maximize density, defined as the fraction of non-NaN cells in the resulting pivot. We do not search over subsets of attribute values, since this would make the pivot search combinatorial and impractical for tables with many tuples.

As shown in Figure 7, even a simple one-level row and column pivot discards 95.3% of the candi-

dates, while requiring at least two hierarchical levels on either rows or columns discards 98.6%. Thus, common non-relational perturbations such as dense pivots and hierarchical headers are rarely directly applicable to real relational sources. The constraint becomes even stronger for parallel and sequential multi-table scenarios, which require compatible evidence values, valid lookup chains, sufficient table sizes, and, when units are perturbed, convertible numerical quantities. By generating the latent source and the perturbed views jointly, GRADINO satisfies these constraints by construction.

Non-relationalization amplifies errors in multi-table settings, even in real tables. Starting from the two relational tables that support at least two hierarchical levels on rows or columns, we extend the source pool with five additional tables from three databases crawled from Kaggle²: Northwind³, Bikestores⁴ and SQL Bootcamp⁵. We then use GRADINO to apply all structural and cell-level perturbations, and generate both single-table instances and 3-table parallel scenarios. This yields 108 single-table and 84 multi-table instances, repeated over three seeds for more stable estimates.

Figure 8 shows that non-relationalization consistently reduces accuracy, but the drop is much larger in the multi-table setting. In single-table instances, performance decreases by 5.7 points for gpt-5-mini, and by 5.5, 7.6, and 13.1 points for Qwen 122B, Qwen 35B, and Qwen 9B. In multi-table instances, the drops increase to 19.8, 25.4, 28.2, and 46.8 points. Importantly, this cannot be explained by the number of tables alone: in the clean relational setting, performance does not decrease when moving from single-table to multi-table inputs. The degradation therefore comes from the interaction between multi-table reasoning and non-relational variability across views.

D Experimental Results and Training Configuration

This appendix provides additional details on the experimental evaluation (§4.4). Specifically, it reports detailed per-setting performance of the evaluated

²used solely for testing, without redistribution

³<https://github.com/microsoft/sql-server-samples/tree/master/samples/databases/northwind-pubs>, MS-PL license

⁴<https://www.sqlservertutorial.net/sql-server-sample-database/>

⁵https://github.com/anshlambagit/SQL_Bootcamp

baselines on parallel and sequential questions, over both relational and non-relational tables generated by GRADINO (§D.1), together with the configuration and results of the finetuning experiments (§D.3).

D.1 Detailed QA Performance

Table 4 reports the performance of gpt-5-mini, Qwen3.5-122B, Qwen3.5-35B, and Qwen3.5-9B on parallel questions in the environmental and financial domains, under both uniform and heterogeneous units of measurement. Consistent with §4.4, accuracy decreases as the number of tables grows (*scaling collapse*) and when compatible quantities are expressed in different units (*unit heterogeneity*).

Table 5 reports the corresponding results on sequential questions. Like in the parallel setting, scaling the amount of tables sharply decreases the performance on sequential questions. Although sequential questions are harder at matched table counts, they are inherently bounded by the number of tables, since realistic lookup chains rarely span more than five tables. Parallel questions, in contrast, can be made arbitrarily harder by increasing the number of tables, as their difficulty is governed by the amount of evidence that must be located and aggregated.

Table 7 reports model performance on questions generated from real tables extracted from Spider and Kaggle. In every setting, questions over non-relational tables are harder than those over their relational counterparts, and this gap is amplified in the multi-table case. While non-relational presentation reduces accuracy by 7.95 points on average in the single-table setting, the drop grows to 30.05 points in the multi-table setting, showing that table count and non-relationality interact to produce a compounded failure mode.

Regarding the use of gpt-5-mini as both the generator and tested model, LLMs are known to favour their own outputs (Panickssery et al., 2024). However, the results show no sign of this: gpt-5-mini is not uniformly ahead of the Qwen models and degrades along the same axes by comparable margins. This is expected, since the properties driving the collapse (table count, perturbations, and unit heterogeneity) are introduced by the rendering stage rather than by the generating model.

D.2 Fine-tuning statistical variance

Table 6 reports the standard deviation of the finetuning results in Table 2. Models finetuned on

Train data	2		3		5		avg
	com	qnt	com	qnt	com	qnt	
None	0.00	0.00	0.00	0.00	0.00	0.00	0.00
GRADINO (fin)	6.18	2.05	4.56	1.60	4.37	1.18	3.32
GRADINO (env)	3.11	3.55	0.59	0.92	2.04	2.04	2.04
TAT-QA	7.40	2.05	0.34	3.70	2.36	1.18	2.84
GRI-QA	3.16	2.71	4.33	2.45	5.57	3.12	3.56
GRADINO (env) + GRI-QA	2.34	1.03	1.89	2.45	1.42	4.08	2.20

Table 6: Standard deviation across 3 seeds of the accuracy reported in Table 2, for Qwen2.5-14B-Instruct finetuned and evaluated on the GRI-QA test split, by number of input tables (2/3/5) and question type (comparative, quantitative). The non-finetuned baseline (None) is deterministic and therefore has zero variance.

GRADINO data exhibit variance comparable to those trained on real data. Moreover, mixing GRADINO data with the GRI-QA training split reduces variance relative to the GRI-QA baseline, suggesting that GRADINO provides not only additional supervision and higher accuracy, but also greater training robustness.

D.3 Training configuration

We finetune Qwen2.5-14B-Instruct (no-thinking) with LoRA (Hu et al., 2022) to assess whether GRADINO-generated data is useful as training supervision (§4.4). LoRA adapters are applied with $r=8$, $\alpha=16$, no dropout and no bias to the query and value attention projections, and the model is loaded in bfloat16 with SDPA attention. Training uses AdamW with a learning rate of 2×10^{-4} , 10 warmup steps, a maximum sequence length of 16,384 tokens, and an effective batch size of 4 (per-device batch size 1 with gradient accumulation 4). We train for 10 epochs under FSDP full-shard data parallelism on 4xA100 GPUs. Following common practice for instruction tuning (Wang et al., 2023), the loss is computed only over the completion tokens, ignoring the prompt. We evaluate at the end of each epoch and retain the checkpoint with the lowest validation loss. All runs are averaged over three different seeds.

Training data. For the GRADINO regimes, training instances are drawn from the generated parallel questions over $\{2, 3, 5\}$ tables, balanced across the three aggregation operators (average, sum, and superlative) and the same-unit and mixed-unit settings. For TAT-QA, we retain only the instances

model	single-table		multi-table		avg
	rel.	non-rel.	rel.	non-rel.	
gpt-5-mini	92.48	86.76	95.24	75.40	87.47
Qwen3.5-122B	93.40	87.93	91.66	66.27	84.81
Qwen3.5-35B	91.88	84.32	92.86	64.68	83.43
Qwen3.5-9B	82.79	69.73	87.30	40.48	70.08
avg	90.14	82.19	91.76	61.71	—

Table 7: Average performance over 3 seeds across models and table settings. Results are grouped into single-table and multi-table (3 tables) settings, with questions over relational and non-relational table views reported separately.

answerable from tables alone, whereas for GRI-QA we partition its multi-table instances into the train and test splits used throughout our evaluation. Across all regimes, we keep only the instances whose answer, computed by gpt-5-mini, was verified as correct, and use the associated reasoning trace as the supervised target, teacher-forcing Qwen2.5-14B-Instruct on the completion only. To ensure a fair comparison across data sources, every regime is capped at the same number of training instances (440), split 90%/10% between training and validation. Models are evaluated on the GRI-QA test split (§4.4, Table 2).

D.4 Program-of-Thoughts results

Table 9 reports the results of gpt-5-mini with Program-of-Thoughts prompting (Chen et al., 2023). When the generated Python code is not executable, we fall back to Chain-of-Thought prompting. Although Program-of-Thoughts yields slightly higher accuracy than Chain-of-Thought in some settings (Table 4), performance still degrades sharply as the number of tables increases and under unit heterogeneity. This suggests that the observed failures are not primarily caused by incorrect arithmetic execution, but by broader difficulties in resolving the question, retrieving the relevant evidence, and handling heterogeneous table presentations.

E Statistics of GRADINO-generated data

Table 8 reports, for each experiment, the number of retained instances, the percentage of generated candidates discarded by the question-verification pass, and how the instances are balanced across table counts, aggregation operators, unit settings, and relational versus non-relational views.

Experiment	Split	# inst.	% disc.	# tables	Question composition (balanced)		
					agg. op.	units	rel./non-rel.
Scaling & units (Fig. 4, 6)	parallel	2,322	22.6	{2, 3, 5, 10, 20}	sum/avg/sup.	same/mixed	—
	sequential	840	6.7	{2, 3, 5}	sum/avg/sup.	same	—
Rel. vs non-rel. (Fig. 8)	single-table	216	0.0	1	sum/avg/sup.	same	✓
	multi-table	168	22.2	3	sum/avg/sup.	same	✓

Table 8: Statistics of the GRADINO-generated data used in our experiments. **# inst.** is the number of instances retained after the question-verification pass (§3); **% disc.** is the share of generated candidates discarded by the schema-format check (§3.1, §B). Under *Question composition*, each split is balanced across the listed table counts, aggregation operators (sum, average, superlative), unit settings (same vs. mixed units, where applicable), and, for the ablation, relational vs. non-relational views. The rel./non-rel. ablation counts both the relational and non-relational rendering of each latent instance ($216 = 108 \times 2$ single-table, $168 = 84 \times 2$ multi-table; cf. §4.4).

Domain	Setting	2			3			5			10			20		
		com	quant	avg	com	quant	avg	com	quant	avg	com	quant	avg	com	quant	avg
Environ.	Same unit	97.37	88.16	92.77	92.11	90.79	91.45	91.89	90.54	91.22	91.67	70.83	81.25	84.85	24.24	54.55
	Unit diff.	60.53	50.00	55.27	50.00	38.16	44.08	40.54	16.22	28.38	22.78	4.17	13.48	24.24	1.52	12.88

Table 9: Accuracy of gpt-5-mini with Program-of-Thoughts prompting (Chen et al., 2023) on the environmental domain by number of input tables and unit setting.

Parallel questions are the most numerous, as they span more table-count settings (including the 10- and 20-table regimes) and are further split into same-unit and mixed-unit variants. They also exhibit the highest discard rate (22.6%). This is a consequence of enforcing the semantic constraints (§B): when the latent source is partitioned into per-table views for parallel questions, constraint enforcement can sharply shrink the feasible table size, leaving some candidate views too small to satisfy the layout requirements (each view is required to have between 3 and 10 rows and 3 and 10 columns). Sequential questions, by contrast, are discarded far less often (6.7%). Their losses stem mainly from hallucinations in the LLM-generated schema, such as malformed JSON or type mismatches (e.g., a categorical attribute declared as float).

The relational-vs-non-relational ablation (Figure 8) behaves differently, because it operates on real tables and therefore skips both schema generation and semantic-constraint enforcement. In the single-table setting no instance is discarded, since the tables are used as-is, although they are themselves the product of an aggressive filtering of Spider, BEAVER, and BIRD (Figure 7). The multi-table setting skips the same two stages but still requires parallel splits: a valid split must isolate the target value by varying a single attribute while keeping each view within the admissible table size

range and pivot density, and such a split cannot always be found. This raises the discard rate to 22.2%, comparable to the synthetic parallel setting despite the absence of constraint enforcement.

F Prompts

This appendix collects the prompts used by GRADINO at each stage of the generation pipeline. Figure 9 shows the prompt used to generate the latent relational source (§3.1), and Figure 10 the prompt used to infer the semantic constraints (§B). Figures 11–13 report the question-generation prompts for the same-unit parallel, mixed-unit parallel, and sequential regimes, while Figures 14–16 report the corresponding question-verification prompts. Finally, Figure 17 shows the inference prompt provided to the evaluated models. All prompts instruct the model to reason in a zero-shot Chain-of-Thought (Kojima et al., 2022) manner.

G Licenses

We indicate the licenses of all the artifacts used. All the artifacts used in this paper can be used for research: Spider (CC BY-SA 4.0), BIRD (CC BY-SA 4.0), BEAVER (MIT), TAT-QA (MIT), GRI-QA (MIT), Qwen models (Apache 2.0).

Relational generation prompt

You are the best table designer in the world for the {domain} topic. You always use lexicon highly specific to {domain}. For the tables you create, you always make the tables "real", using real entities while avoiding placeholders. You always use lexicon that is specific to {domain}, and make textual entries in the tables be composed even by more words (unlike standard relational tables), like in the following examples (you must come up with different lexicon than the one used in the examples):

{examples}

You must use variations of the terminology used previously, but it's important you avoid using highly common terms. Use terms that only a real {domain} expert would know. Then, create the following Python dictionary, where the number of attributes (number of columns) must be exactly equal to num_columns.

```
{
  "name": ..., # name of the table, must be in pascal casing
  "attributes": ..., # list containing the names of the attributes. The names must be written
  ↳ in pascal casing. They must not have whitespaces
  "attributes_long": ..., # list containing the names of the attributes. This list must be
  ↳ exactly the same as the "attributes" list, but the attribute names must not be in camel
  ↳ casing, pascal casing or snake casing. They may contain consist of multiple words. Use
  ↳ {domain}-specific lexicon, as the one used in the "examples" below, but be creative and
  ↳ vary the lexicon. Not every word must have the initial letter in uppercase.
  "attribute_types": ..., \# list containing the types of the attributes. the length of this
  ↳ list must be equal to the length of the "attributes" list. The type must be
  ↳ "categorical", or either "int" or "float" for the "value_col" column. Also numerical
  ↳ values (like years or IDs) can be categorical. Only the value_col column can be float.
  "range": ..., # this list must have the same length as "attributes" and "attribute_types".
  ↳ For categorical attributes it is a list containing {col_cardinality} different values,
  ↳ which can be lengthy like in standard web tables. Use {domain}-specific lexicon like
  ↳ the one used in the examples below, but be creative and vary the lexicon. For possible
  ↳ float and int values it is a list containing, at the first position, the start of the
  ↳ range and at the second position the end of the range (extremes included).
  "value_col": ..., # string indicating the attribute name of the column to pivot later. The
  ↳ name of the attribute must be one of the names in "attributes". This table attribute
  ↳ must contain values that are either "int" or "float".
  "unit_of_measurement": ..., # one string, where the unit must be one of the following
  ↳ units: {units}
  "number_of_decimals": ... # integer value indicating the number of decimals (integer
  ↳ greater or equal than 0) to use when representing values in the "value_col" column.
}
```

The generated table must NOT use the following attributes and values: Attributes of past tables: {past} Corresponding values of past tables: {past_values}

First reason step-by-step. Then write exactly "Final answer: " (the text must be exactly the same) followed exclusively by the requested Python dictionary. To avoid numerical inconsistencies, you must not specify any type of unit of measurement inside the table ("name", "attributes", "attributes_long", "attribute_types", "range", "value_col"), but only inside the "unit_of_measurement" field. This is very important. Ensure the output is in the expected format. Make sure that the proposed table is about {domain} and at the same time does uses completely different attributes and values as the tables used previously. Make sure that the table values ("range") do not include any specification about units of measurement, as the unit of measurement must be specified only in the "unit_of_measurement" field. This is very important to avoid numerical inconsistencies. Make sure to write exactly "Final answer: " at the end (the text, together with ":", must be exactly and completely the same), followed by the required dictionary. Do not write anything else after "Final answer: ".

Let's think step-by-step.

Figure 9: Prompt used by gpt-5-mini to generate the latent relational source.

Semantic constraint generation prompt

You are an Operations Research expert in the {domain} domain. Given table metadata, infer realistic constraints among attributes and values.

Input

You will receive a JSON-like object:

{schema_json}

Task

Infer constraints using domain knowledge. Produce:

- inter-row constraints: relationships between *different rows*, always comparing the numeric value in 'value_col'

- intra-row constraints: relationships between *attributes within the same row*

Do not create constraints that are already explicitly defined by the table structure (e.g., bounds in the range list, unit of measurement, number of decimals...).

Row selection syntax

- Select a row by fixing categorical values: Attribute == "Value" (the operator can be a python boolean operator)

- Combine selectors: (Attribute1 == "Value1" and Attribute2 == "Value2" and ...)

- Refer to the numeric value using the 'value_col' name as a property: (...selectors...).{{value_col}} where {{value_col}} is the placeholder containing the value_col attribute name, not the string "value_col", and must not be used with the {{{}} (which are part of the placeholder)

Example selectors:

- (Company == "Ferrari")

- (Year == 2020 and TypeOfMoney == "Gross") String values expect to be enclosed within "", while numerical data must not.

Constraint language

Write each rule as a Python-valid boolean/math expression (>, >=, ==, !=, +, -, *, /, and, or, not, parentheses).

- Inter-row constraints MUST compare '(...selector...).{{value_col}}' between two (or more) selected rows. The selector must be surrounded by brackets. Across each selected row, all non-specified attributes are considered to have the same value.

- Intra-row constraints MUST compare attributes in the same row and must use conditional form. The selector and expression must be surrounded by brackets:

"if (...selector...) then (<expression>)"

Examples (inter-row):

- (TypeOfMoney == "Gross").{{value_col}} > (TypeOfMoney == "Net").{{value_col}}

- (Year == 2020 and TypeOfMoney == "BudgetYearBeginning").{{value_col}} == (Year == 2019 and TypeOfMoney == "BudgetYearEnd").{{value_col}} (this applies to the same Company in both rows, as Company is not specified in the selectors)

Examples (intra-row):

- if (CompanyTier == "Small") then ({{value_col}} <= 5000000)

- if (TypeOfMoney == "COGS" or Scenario == "Actual") then (Scenario != "Forecast") The expression (e.g. Scenario != "Forecast") cannot contain "and", "or" or "not" operators (for the "and", multiple rules must be created, instead of using "and")

When you have lhs or rhs with mathematical operations (not boolean operations), enclose the lhs or rhs with "()" parentheses.

Output format (strict)

First, reason step-by-step. Then output exactly "Final answer:" followed by a Python dictionary in this format:

```
{
  "inter_row_constraints": [<string>, ...], # the list of strings can be empty if no inter-row
    ↪ constraints are found
  "intra_row_constraints": [<string>, ...] # the list of strings can be empty if no intra-row
    ↪ constraints are found
}
```

Make sure that you use the exact attribute names and values as provided in the metadata, and that you exactly follow the constraint language and output format. Remember that {{value_col}} is the placeholder containing the value_col attribute name, not the string "value_col", and must not be used with the {{{}} (which are part of the placeholder)

Table info: {input}

Let's think step-by-step.

Figure 10: Prompt used by gpt-5-mini to generate the semantic constraints.

Parallel same unit question generation prompt

Given a list of SQL queries, a list of HTML tables, a final aggregation operator across the tables, and the final result, you must generate a natural language question. In particular, to obtain the final result, each SQL query is executed on its corresponding table, and then the aggregation operator is applied to the intermediate results. The natural language question must be answerable by the same result. However, you must differ how the sentence is phrased and the words used, in order to avoid too much overlap between question and table contents. You must not use the exact SQL table attributes, where the attribute name is in camel case, but you must use them in a more natural way, resembling user questions. First reason step-by-step. Then, write exactly "Final question:" followed exclusively by the novel natural language question. Make sure to include all the information needed to answer correctly the question, so that the question is still answerable with the original SQL result. Do not write anything else after "Final question:".

{text}

The generated question must be of the following type: {method}

For superlative questions, make sure to ask for the value. Never ask for the entity.

The final aggregation operation, that is applied to the extracted values from each table, is the following: {aggregation}

This is the final result: {result}

Make sure that the generated question also explicits the final aggregation operation applied over the extracted values. The question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). Also, the natural language question must not miss any crucial information needed to answer the question. Let's think step-by-step.

Figure 11: Prompt used by gpt-5-mini to generate the natural language question for same-unit parallel instances.

Parallel mixed unit question generation prompt

Given a list of SQL queries, a list of HTML tables, a final aggregation operator across the tables, an unit of measurement and the final result, you must generate a natural language question. In particular, to obtain the final result, each SQL query is executed on its corresponding table, the values must be normalized to the requested unit of measurement, and then the aggregation operator is applied to the intermediate results. The natural language question must be answerable by the same result. However, you must differ how the sentence is phrased and the words used, in order to avoid too much overlap between question and table contents. You must not use the exact SQL table attributes, where the attribute name is in camel case, but you must use them in a more natural way, resembling user questions. First reason step-by-step. Then, write exactly "Final question:" followed exclusively by the novel natural language question. Make sure to include all the information needed to answer correctly the question, so that the question is still answerable with the original SQL result. Do not write anything else after "Final question:".

{text}

The generated question must be of the following type: {method}

For superlative questions, make sure to ask for the value. Never ask for the entity.

The final aggregation operation, that is applied to the extracted values from each table, is the following: {aggregation}

This is the final result: {result}

Make sure that the generated question also explicits the final aggregation operation applied over the extracted values. Make sure that the final natural language question always specifies the following unit of measurement: {unit} The question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). Also, the natural language question must not miss any crucial information needed to answer the question. Let's think step-by-step.

Figure 12: Prompt used by gpt-5-mini to generate the natural language question for mixed-unit parallel instances.

Sequential question generation prompt

You will be given some SQL queries that are applied sequentially to each table, some tables and the result of executing those SQL queries. You must generate a single natural language question that answers the final result. The question must be phrased exactly with the following style:

"What is the [operation] [attribute_value with unit of measure] for [constraints]?"

Restrict the calculation to:

- 1) first option in last SQL query
- 2) second option in last SQL query
- ..."

In particular,

- (1) the [constraints] must be the WHERE constraints appearing inside the first SQL query;
- (2) the options must be the options appearing inside the last SQL query;
- (3) by looking at the tables' contents, make sure that the question is semantically sound and plausible;
- (4) you must differ how the sentence is phrased and the words used, in order to avoid too much overlap between question and table contents, as well as overlap with the value of the attributes;
- (5) you must not use the exact SQL table attributes, where the attribute name is in camel case, but you must use them in a more natural way, resembling user questions;
- (6) when possible, prefer using shorter questions: the less tokens you use the better, without losing any important details, such as the constraints or the final question. Variations in the formulation of the values of the attributes between NL and SQL questions are advised, as long as the question remains unambiguous.

Additionally:

- (1) the generated question must be of the following type: {method}
- (2) for superlative questions, make sure to ask for the value. Never ask for the entity.

First reason step-by-step. In the end, write "Final question:" followed exclusively by the final question. Do not write anything else after "Final question:".

{text}

This is the final result: {result}

Let's think step-by-step.

Figure 13: Prompt used by gpt-5-mini to generate the natural language questions for sequential instances.

Parallel same unit question verification prompt

You will be given a question in natural language, some SQL extraction queries, some tables and the result of executing those SQL queries + an aggregation operation (e.g. sum, average, max/min extraction or similar) on those tables. Your task is to verify if the question in natural language is equivalent to the aggregation of the SQL extractive queries and produces the same result. You must check if the natural language question does not specify a needed constraint that the SQL queries specifies. Also, you must check if the natural language question does not specify the final aggregation operation to apply. First, reason step-by-step and analyze the natural language question and the SQL queries + aggregation operation to determine if they are asking for the same information. If they are equivalent (use the same information), in the end write exactly "Final answer: Yes". If they are not equivalent, in the end write exactly "Final answer:" followed exclusively by the novel formulation of the natural language question that would make it equivalent to the SQL queries + aggregation operation. In this case, the natural language question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). Do not write anything else after "Final answer:".

Natural Language Question: {nl_question}

SQL Question: {sql_question}

Tables: {table}

Result: {sql_result}

Remember that, if you change the question, the natural language question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine).

Let's think step-by-step.

Figure 14: Prompt used by gpt-5-mini to verify the natural language question for same-unit parallel instances.

Parallel mixed unit question verification prompt

You will be given a question in natural language, some SQL extraction queries, some tables and the result of executing those SQL queries + an aggregation operation (e.g. sum, average, max/min extraction or similar) on those tables. Also, you will be provided with the reference unit of measurement to use inside the question. Your task is to verify if the question in natural language is equivalent to the aggregation of the SQL extractive queries and produces the same result, considering also possible unit of measurement normalizations. You must check if the natural language question does not specify a needed constraint that the SQL queries specifies. Also, you must check if the natural language question does not specify the final aggregation operation to apply or the final unit of measurement the answer expects. First, reason step-by-step and analyze the natural language question and the SQL queries + aggregation operation + unit of measurement to determine if they are asking for the same information. If they are equivalent (use the same information), in the end write exactly "Final answer: Yes". If they are not equivalent, in the end write exactly "Final answer:" followed exclusively by the novel formulation of the natural language question that would make it equivalent to the SQL queries + aggregation operation + unit of measurement. In this case, the natural language question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). The final natural language question must always specify the following unit of measurement: {unit}. Do not write anything else after "Final answer:".

Natural Language Question: {nl_question}

SQL Question: {sql_question}

Tables: {table}

Result: {sql_result}

Remember that, if you change the question, the natural language question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). Also, the final natural language question must always specify the following unit of measurement: {unit}.

Let's think step-by-step.

Figure 15: Prompt used by gpt-5-mini to verify the natural language question for mixed-unit parallel instances.

Sequential question verification prompt

You will be given a question in natural language, SQL queries that are applied sequentially, some tables and the result of executing those SQL queries. You must check if the natural language question does not specify a needed constraint that, without it, makes it impossible to correctly answer the question. If there are any missing constraint, your MUST explicitly add those constraints, without unnecessarily modifying other parts of the question.

The question must be formulated exactly with the following style:

"What is the [operation] [attribute_value with unit of measure] for [constraints]?"

Restrict the calculation to:

- 1) first option in last SQL query
- 2) second option in last SQL query
- ..."

First, write your step-by-step reasoning and analyze the natural language question and the SQL queries to determine if there are any missing constraints. If they are equivalent (use the same information), at the end of your reasoning write exactly "Final answer: Yes". If they are not equivalent, at the end of your reasoning write exactly "Final answer:" followed exclusively by the novel formulation of the question with the missing constraints. Do not write anything else after "Final answer:". Slight variations in the formulation of the values of the attributes between NL and SQL questions are advised, as long as the question remains unambiguous.

Natural Language Question: {nl_question}

SQL Question: {sql_question}

Tables: {table}

Result: {sql_result}

Remember that, if you change the question, the natural language question must be user-like: it must not contain SQL syntax, camel case, mathematical or boolean symbols, or explicit mentions of certain multi-word string values (appearing in the table) enclosed in "" (slight variations are fine). The final question must always clearly explicitate what are the values that are needed to apply the aggregation operation in the end. It must not say "what is the value across the specified indicators", but say something among the lines of "what is the value across: (1) ..., (n) ...". This is important, as the user cannot see the SQL query, and everything needs to be explicitly stated. Also, the final natural language question must always specify the following unit of measurement: {unit} Remember also to first reason step-by-step. In the end, write "Final answer:" followed exclusively by the answer. Do not write anything else after "Final answer:". Let's think step-by-step.

Figure 16: Prompt used by gpt-5-mini to verify the natural language question for sequential instances.

Inference prompt

Answer the following question given the provided HTML table.

First reason step-by-step, then write "Final answer:" followed exclusively by the correct answer. Do not write anything else after "Final answer:"

Every calculation must be done with a precision of exactly 6 decimal places.

Only the numerical value must be written in the final answer.

Question: {question}

Table: {table}

Let's think step-by-step.

Figure 17: Inference prompt. Even though the EM metric requires 2-decimal precision (§4.2), the prompt asks for a broad numerical precision of 6 decimals values to prevent premature truncation during the model's intermediate computation, which could otherwise accumulate small rounding errors across successive operations.